

MHT Simulation

Connor Demorest


9/4/2021

Simulation study:

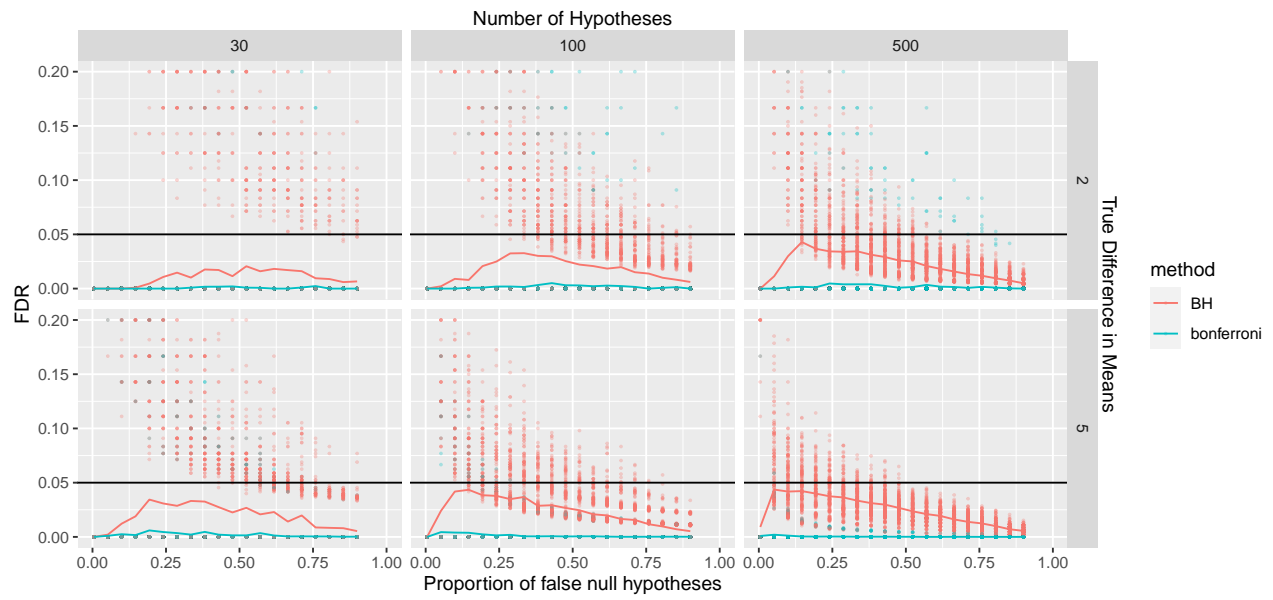
- 1) Take 100 samples from $N(0,1)$ and $N(5,1)$ mixed with proportion p in $(0,1)$. Take $\alpha = 0.05$. $X_i = (1 - p) * N(0, 1) + p * N(3, 1)$
- 2) Find p values of X
- 3) Use BH, etc
- 4) Get FDP and Power
- 5) Repeat many times to find FDR and avg power across values of p

As a function:

Power of variations of parameters for FDR control methods

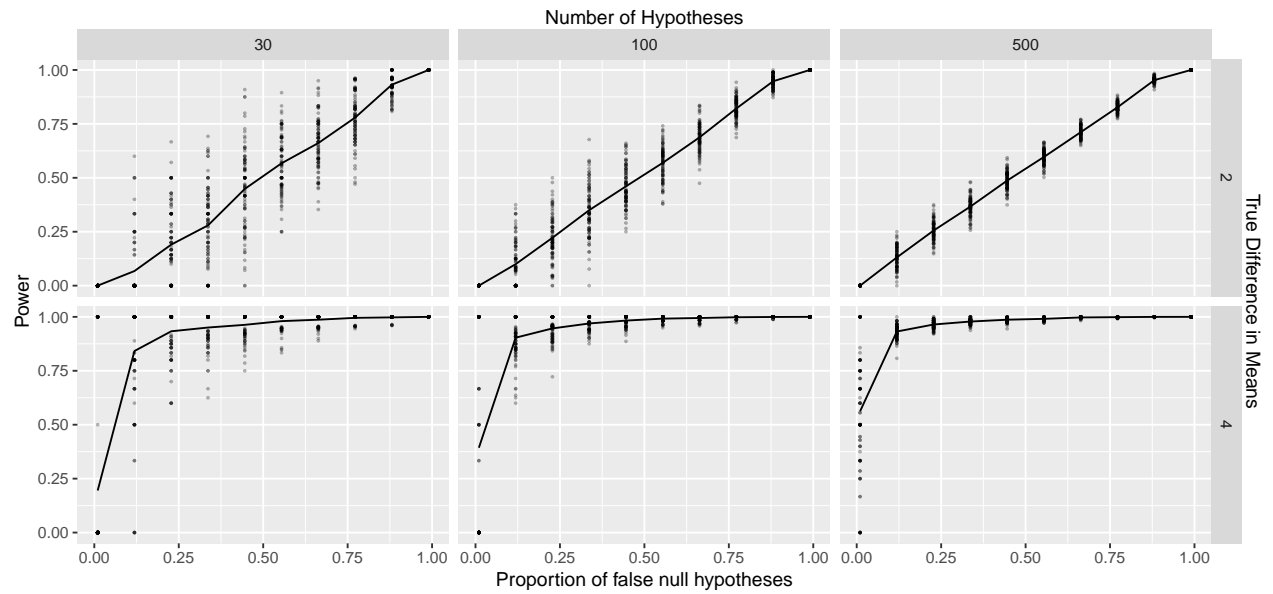


FDR for variations of parameters for FDR control methods

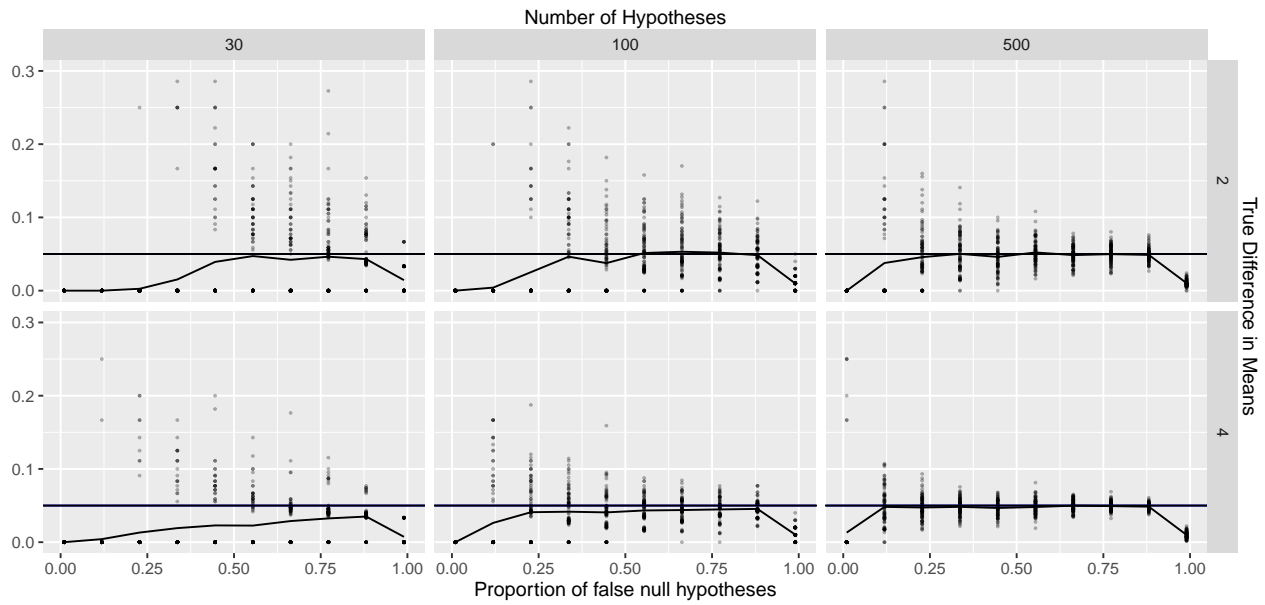


Using Lfdr and PFDR

Lfdr control method simulation



Lfdr control method simulation



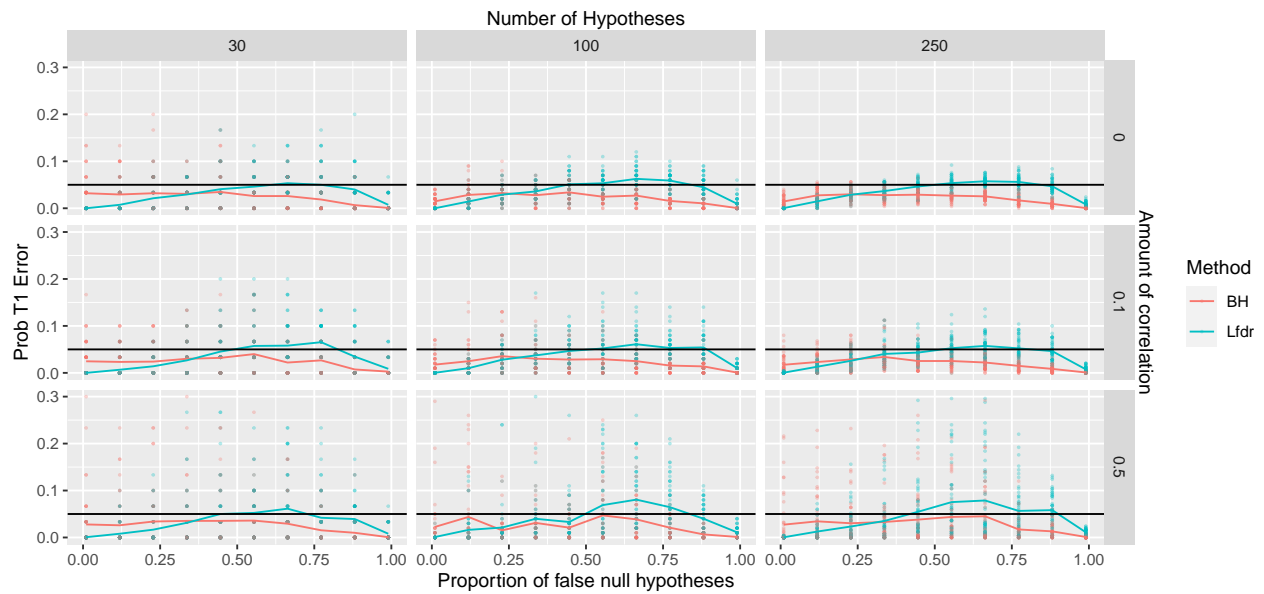
Using the same setting as the Lfdr procedure, compare the performance of the Lfdr and the BH procedure,

- When the variances are same, i.e., 1 in both null and non-null groups -> *see above*
- Draw the samples from a multivariate normal distribution. For example, if you have 100 samples, out of which $\pi = 0.5$ (say), proportion are from the non-null distribution, draw one sample from the following: Define $\mu_vector = (1-\theta)f(0) + (\theta)f(2)$, i.e. a 100x1 vector of 0 (null mean) and 2(non-null mean). $\Sigma =$ a 100x100 matrix, with diagonal elements 1 and off diagonal elements = ρ (a positive value between 0 and 1). Draw one sample from $MVNorm_100(\mu = \mu_vector, \sigma = \Sigma)$ Find the p-values, and Apply the BH and Lfdr process on this sample to figure out the discoveries, false and true discoveries. This would show the performance of the two procedures when the p-values are positively dependent.


```
## Warning: Removed 19 rows containing non-finite values (stat_summary).
```

```
## Warning: Removed 19 rows containing missing values (geom_point).
```

FDR control methods in positive dependence



FDR control methods in positive dependence



Should I have been doing two-sided tests?

Code:

```
knitr::opts_chunk$set(echo = F, fig.width = 10, fig.height = 5,
                       cache = TRUE)

library(tidyverse)
p = 0.4
a = 0.05
num_samples = 100
# True if true mean != 0
A = rbernoulli(num_samples, p)
t_nulls = rnorm(num_samples, 0, 1)
```

```

f_nulls = rnorm(num_samples, 2, 1)
samples = ifelse(A, f_nulls, t_nulls)
p_vals = pnorm(samples, lower.tail = F)

# Test discoveries
B = (p_vals <= a)
# FP = A' & B
## False_positive = (A & B)
# FDP = P(A' & B)/P(B)
FDP = ifelse(is.na(sum(!A & B)/sum(B)), 0, sum(!A & B)/sum(B))

# Power = P(reject Null | Null false) = P(reject & Null false)/P(Null false)
Power = ifelse(is.na(sum(A & B)/sum(A)), 0, sum(A & B)/sum(A))
# m === number of total hypotheses
func = function(p = 0.4, a = 0.05, m = 1000, alt_hyp = 3, method = "none") {
  # A === True discoveries
  A = rbernoulli(m, p)
  t_nulls = rnorm(m, 0, 1)
  f_nulls = rnorm(m, alt_hyp, 1)
  samples = ifelse(A, f_nulls, t_nulls)
  p_vals = pnorm(samples, lower.tail = F)
  adj_p_vals = p.adjust(p = p_vals, method = method)

  # Test discoveries
  B = (adj_p_vals <= a)
  # FP = A' & B
  False_positive = ifelse(!A & B, 1, 0)
  # FDP = P(A' & B)/P(B)
  FDP = ifelse(is.na(sum(!A & B)/sum(B)), 0, sum(!A & B)/sum(B))

  # Power = P(reject Null | Null false) = P(reject & Null false)/P(Null false)
  Power = ifelse(is.na(sum(A & B)/sum(A)), 0, sum(A & B)/sum(A))
  return(c(FDP, Power))
}

# New function: do_many, where it just does 'many' and returns averages?
# Calculate FDR and Avg Power

# many = replicate(1000, func(method = "none")) %>% apply(., 1, FUN = mean) %>% print

df = expand_grid(p = rep(seq(0.005, 0.9, length.out = 20), each = 200),
  method = c("bonferroni", "BH"),
  alt_hyp = c(2, 5),
  m = c(30, 100, 500))
many = t(mapply(func,
  p = df$p,
  method = df$method,
  alt_hyp = df$alt_hyp,
  m = df$m)) %>%
  as_tibble %>%
  rename(FDR = V1, Power = V2) %>%
  mutate(df)

# do_many = function(num_reps = 1000, p = 0.4, a = 0.05, num_samples = 1000, method = "none") {

```

```

# many = replicate(num_reps, func(p = p, a = a, num_samples = num_samples, method = method)) %>%
#   apply(., 1, FUN = mean)
#   return(many)
# }
#
# df = expand_grid(p = seq(0.1, 0.9, by = 0.01), method = c("none", "BH"))
# df2 = t(mapply(do_many, p = df$p, method = df$method, num_reps = 100)) %>%
#   as_tibble %>%
#   rename(FDR = V1, Power = V2) %>%
#   mutate(df)

# Increased p gives more true discoveries
ggplot(data = many, aes(x = p, y = Power, color = method)) +
  geom_point(size = 0.3, alpha = 0.2) +
  stat_summary(fun = "mean", geom = "line") +
  facet_grid(alt_hyp ~ m) +
  labs(x = "Proportion of false null hypotheses",
       title = "Power of variations of parameters for FDR control methods") +
  scale_y_continuous(sec.axis = sec_axis(~ ., name = "True Difference in Means",
                                         breaks = NULL, labels = NULL), limits = c(0,1)) +
  scale_x_continuous(sec.axis = sec_axis(~ ., name = "Number of Hypotheses",
                                         breaks = NULL, labels = NULL), limits = c(0,1))

ggplot(data = many, aes(x = p, y = FDR, color = method)) +
  geom_point(size = 0.3, alpha = 0.3) +
  stat_summary(fun = "mean", geom = "line") +
  facet_grid(alt_hyp ~ m) +
  geom_abline(intercept = 0.05, slope = 0) +
  labs(x = "Proportion of false null hypotheses",
       title = "FDR for variations of parameters for FDR control methods") +
  scale_y_continuous(sec.axis = sec_axis(~ ., name = "True Difference in Means",
                                         breaks = NULL, labels = NULL), limits = c(0,0.2)) +
  scale_x_continuous(sec.axis = sec_axis(~ ., name = "Number of Hypotheses",
                                         breaks = NULL, labels = NULL), limits = c(0,1))

# Basically re-using code from above
# p known here!
p = 0.4
alpha = 0.05
num_samples = 1000
## True if true mean != 0
# theta = rbernoulli(num_samples, p)
# f0 = rnorm(num_samples, 0, 1)
# f1 = rnorm(num_samples, 5, 1)
# X = ifelse(theta, f1, f0)
# f = (1-p)*dnorm(X, 0, 1) + p*dnorm(X, 5, 1)

# Lfdr = (1-p)f0(X)/f(x)
PFDR_fun = function(p = 0.4, alpha = 0.05, num_samples = 1000, alt_hyp = 4) {
  df = tibble(theta = rbernoulli(num_samples, p),
              X = ifelse(theta, rnorm(num_samples, alt_hyp, 1), rnorm(num_samples, 0, 1)),
              f = (1-p)*dnorm(X, 0, 1) + p*dnorm(X, alt_hyp, 1),
              Lfdr = (1-p)*dnorm(X, 0, 1)/f) %>%
  arrange(Lfdr) %>%

```

```

mutate(q = cumsum(Lfdr)/(1:length(Lfdr)),
      `q<a` = q < alpha, # Total discoveries
      Type1error = `q<a`&!theta, # Type 1 errors
      Type2error = theta&!`q<a`) %>% # Type 2 errors
      summarise(theta = sum(theta), disc = sum(`q<a`), t1 = sum(Type1error), t2 = sum(Type2error), p, alt,
      return(df)
}

# PFDR_fun()
# reps = plyr::rdply(100, PFDR_fun(), .id = NULL) %>% summarise(PFDR = t1/disc)
#
# mean(as.matrix(reps))

#### Old code ####
# PFDR = rep(NA, length(Lfdr))
# Q = rep(NA, length(Lfdr))
# for(t in 1:length(Lfdr)) {
#   Q[t] = (sum(sort(Lfdr)[1:t]/t))# %>% print
#   PFDR[t] = ((sum(sort(Lfdr)[1:t]/t)) < alpha)# %>% print
# }
####

pr = seq(0.01, 0.99, length.out = 10)
alt = c(2, 4)
samps = c(30, 100, 500)
pr_alt_samps = expand.grid(pr, alt, samps)

# If no discoveries, then no type 1 errors and fdr = 0
reps = replicate(100,
  mapapply(PFDR_fun,
    p = pr_alt_samps[,1],
    alt_hyp = pr_alt_samps[,2],
    num_samples = pr_alt_samps[,3]) %>% t(),
    simplify = F) %>%
do.call(rbind, .) %>%
as_tibble() %>%
mutate(across(.col = everything(), as.double),
      PFDR = ifelse(is.na(t1/disc), 0, t1/disc),
      Power = 1 - ifelse(is.na(t2/theta), 1, t2/theta))

# ### Make plots! Change values of p, 10 are plenty
# ggplot(reps, mapping = aes(x = p, y = PFDR)) +
#   geom_point(size = 0.3, alpha = 0.3, mapping = aes(color = "FDR")) +
#   geom_hline(yintercept = alpha, color = "blue") +
#   stat_summary(fun = "mean", geom = "line", mapping = aes(color = "FDR")) +
#   geom_point(size = 0.3, alpha = 0.3,
#     mapping = aes(x = p, y = Power, color = "Power")) +
#   stat_summary(fun = "mean", geom = "line",
#     mapping = aes(x = p, y = Power, color = "Power")) +
#   facet_grid(alt_hyp ~ num_samples) +
#   geom_abline(intercept = 0.05, slope = 0) +
#   labs(x = "Proportion of false null hypotheses",
#     title = "Lfdr control method simulation",
#     y = "",

```

```

#       color = "") +
#   scale_y_continuous(sec.axis = sec_axis(~ . , name = "True Difference in Means",
#                                           breaks = NULL, labels = NULL), limits = c(0,1)) +
#   scale_x_continuous(sec.axis = sec_axis(~ . , name = "Number of Hypotheses",
#                                           breaks = NULL, labels = NULL), limits = c(0,1))

ggplot(reps, mapping = aes(x = p, y = Power)) +
  # geom_point(size = 0.3, alpha = 0.3, mapping = aes(color = "FDR")) +
  # geom_hline(yintercept = alpha, color = "blue") +
  # stat_summary(fun = "mean", geom = "line", mapping = aes(color = "FDR")) +
  geom_point(size = 0.3, alpha = 0.3) +
  stat_summary(fun = "mean", geom = "line") +
  facet_grid(alt_hyp ~ num_samples) +
  labs(x = "Proportion of false null hypotheses",
       title = "Lfdr control method simulation",
       y = "Power") +
  scale_y_continuous(sec.axis = sec_axis(~ . , name = "True Difference in Means",
                                           breaks = NULL, labels = NULL), limits = c(0,1)) +
  scale_x_continuous(sec.axis = sec_axis(~ . , name = "Number of Hypotheses",
                                           breaks = NULL, labels = NULL), limits = c(0,1))

### Make plots! Change values of p, 10 are plenty
ggplot(reps, mapping = aes(x = p, y = PFDR)) +
  geom_point(size = 0.3, alpha = 0.3) +
  geom_hline(yintercept = alpha, color = "blue") +
  stat_summary(fun = "mean", geom = "line") +
  facet_grid(alt_hyp ~ num_samples) +
  geom_abline(intercept = 0.05, slope = 0) +
  labs(x = "Proportion of false null hypotheses",
       title = "Lfdr control method simulation",
       y = "",
       color = "") +
  scale_y_continuous(sec.axis = sec_axis(~ . , name = "True Difference in Means",
                                           breaks = NULL, labels = NULL), limits = c(0,0.3)) +
  scale_x_continuous(sec.axis = sec_axis(~ . , name = "Number of Hypotheses",
                                           breaks = NULL, labels = NULL), limits = c(0,1))

# p = 0.4
# num_samples = 100
# theta = rbernoulli(num_samples, p)
# f0 = rnorm(num_samples, 0, 1)
# f1 = rnorm(num_samples, 2, 1)
# mu_vec = ifelse(theta, f1, f0)
# rho = 0.2
# Sigma = matrix(nrow = num_samples, ncol = num_samples, data = rho)
# diag(Sigma) = 1
#
# MVN_100 = MASS::mvrnorm(n = 1, mu = mu_vec, Sigma = Sigma)
# MVN_p = pnorm(MVN_100, mean = 0, sd = 1, lower.tail = F)
# BH_p = p.adjust(MVN_p, method = "BH")
#
# df = tibble(theta, BH_p,
#             mu_vec,
#             f = (1-p)*dnorm(mu_vec, 0, 1) + p*dnorm(mu_vec, 2, 1),

```



```

#           Lfdr = (1-p)*dnorm(mu_vec, 0, 1)/f) %>%
#   arrange(Lfdr) %>%
#   mutate(q = cumsum(Lfdr)/(1:length(Lfdr)),
#           `q<a` = q < alpha, # Total discoveries
#           Type1error = `q<a`&!theta, # Type 1 errors
#           Type2error = theta&`!q<a`)

# Takes in rho, N, proportions, and the alternative hypothesis
# Outputs the T1 and T2 error rates for BH and Lfdr
BH_lfdr_fun = function(num_samples = 100, p = 0.4, rho = 0, alt_hyp = 3) {
  # Sigma has correlation structure rho and 1 on diagonal
  Sigma = matrix(nrow = num_samples, ncol = num_samples, data = rho)
  diag(Sigma) = 1
  # For known p, simulate the true hypothesis state
  theta = rbernoulli(num_samples, p)
  # Maybe I don't need these next three lines?
  f0 = rnorm(num_samples, 0, 1)
  f1 = rnorm(num_samples, alt_hyp, 1)
  mu_vec = ifelse(theta, f1, f0)
  # Generate positively correlated data
  MVN = MASS::mvrnorm(n = 1, mu = mu_vec, Sigma = Sigma)
  # Before, I was using the wrong values for MVN below here! Eek!
  # Get p-values for BH and adjust them
  MVN_p = pnorm(MVN, mean = 0, sd = 1, lower.tail = F)
  BH_p = p.adjust(MVN_p, method = "BH")
  # Copy/paste code from before to get Lfdr
  df = tibble(theta,
              BH_p,
              f = (1-p)*dnorm(MVN, 0, 1) + p*dnorm(MVN, alt_hyp, 1),
              Lfdr = (1-p)*dnorm(MVN, 0, 1)/f) %>%
  arrange(Lfdr) %>%
  mutate(q = cumsum(Lfdr)/(1:length(Lfdr)),
         `q<a` = q < 0.05,
         # Calculate errors for BH and Lfdr methods
         Lfdr_Type1error = `q<a`&!theta,
         Lfdr_Type2error = theta&`!q<a`,
         BH_Type1error = (BH_p < 0.05) & !theta,
         BH_Type2error = theta & !(BH_p < 0.05))
  # Return df of mean for Lfdr and BH error rates
  # Fancy version of apply if you're not exposed to map before
  return(map_df(df[, 7:10], mean))
}

df = expand.grid(p = rep(seq(0.01, 0.99, length.out = 10), times = 50),
               num_samples = rep(c(30, 100, 250), times = 1),
               rho = rep(c(0, 0.1, 0.5), times = 1))

# Slow as heck without running in parallel
parallel::mcmapply(FUN = BH_lfdr_fun,
                  mc.cores = parallel::detectCores(),
                  p = df$p,
                  num_samples = df$num_samples,
                  rho = df$rho,

```

```

    alt_hyp = 2) -> a

# Move data around into a nice way to plot using ggplot
BH_Lfdr = t(a) %>%
  as_tibble(column_name = rownames(a)) %>%
  cbind(df) %>%
  mutate_all(as.double) %>%
  pivot_longer(cols = ends_with("error"),
               names_to = c("Method", "Error"),
               names_sep = "_",
               values_to = "Prob") %>%
  mutate(Error = str_sub(Error, 1, 5))

# Plot things
ggplot(BH_Lfdr %>% filter(Error == "Type1"),
       mapping = aes(x = p, y = Prob, color = Method)) +
  geom_point(size = 0.3, alpha = 0.3) +
  stat_summary(fun = "mean", geom = "line") +
  geom_hline(yintercept = 0.05) +
  facet_grid(rho ~ num_samples) +
  labs(x = "Proportion of false null hypotheses",
       title = "FDR control methods in positive dependence",
       y = "Prob T1 Error") +
  scale_y_continuous(sec.axis = sec_axis(~ ., name = "Amount of correlation",
                                           breaks = NULL, labels = NULL), limits = c(0,0.3)) +
  scale_x_continuous(sec.axis = sec_axis(~ ., name = "Number of Hypotheses",
                                           breaks = NULL, labels = NULL), limits = c(0,1))

ggplot(BH_Lfdr %>% filter(Error == "Type2"),
       mapping = aes(x = p, y = 1-Prob, color = Method)) +
  geom_point(size = 0.3, alpha = 0.3) +
  stat_summary(fun = "mean", geom = "line") +
  facet_grid(rho ~ num_samples) +
  labs(x = "Proportion of false null hypotheses",
       title = "FDR control methods in positive dependence",
       y = "Power") +
  scale_y_continuous(sec.axis = sec_axis(~ ., name = "Amount of correlation",
                                           breaks = NULL, labels = NULL), limits = c(0,1)) +
  scale_x_continuous(sec.axis = sec_axis(~ ., name = "Number of Hypotheses",
                                           breaks = NULL, labels = NULL), limits = c(0,1))

```